



zoom AI-Contact Center

Meet The Team



Helen Vu

Ozioma Anele

Wyatt Boska

Hemnathe Gopal

Rinad Hussein

Preview

1. Vision
2. Market Overview & Customer Insights
3. Competitive Analysis
4. Product Strategy, Features & Roadmap
5. Prototype
6. Go To Market Strategy

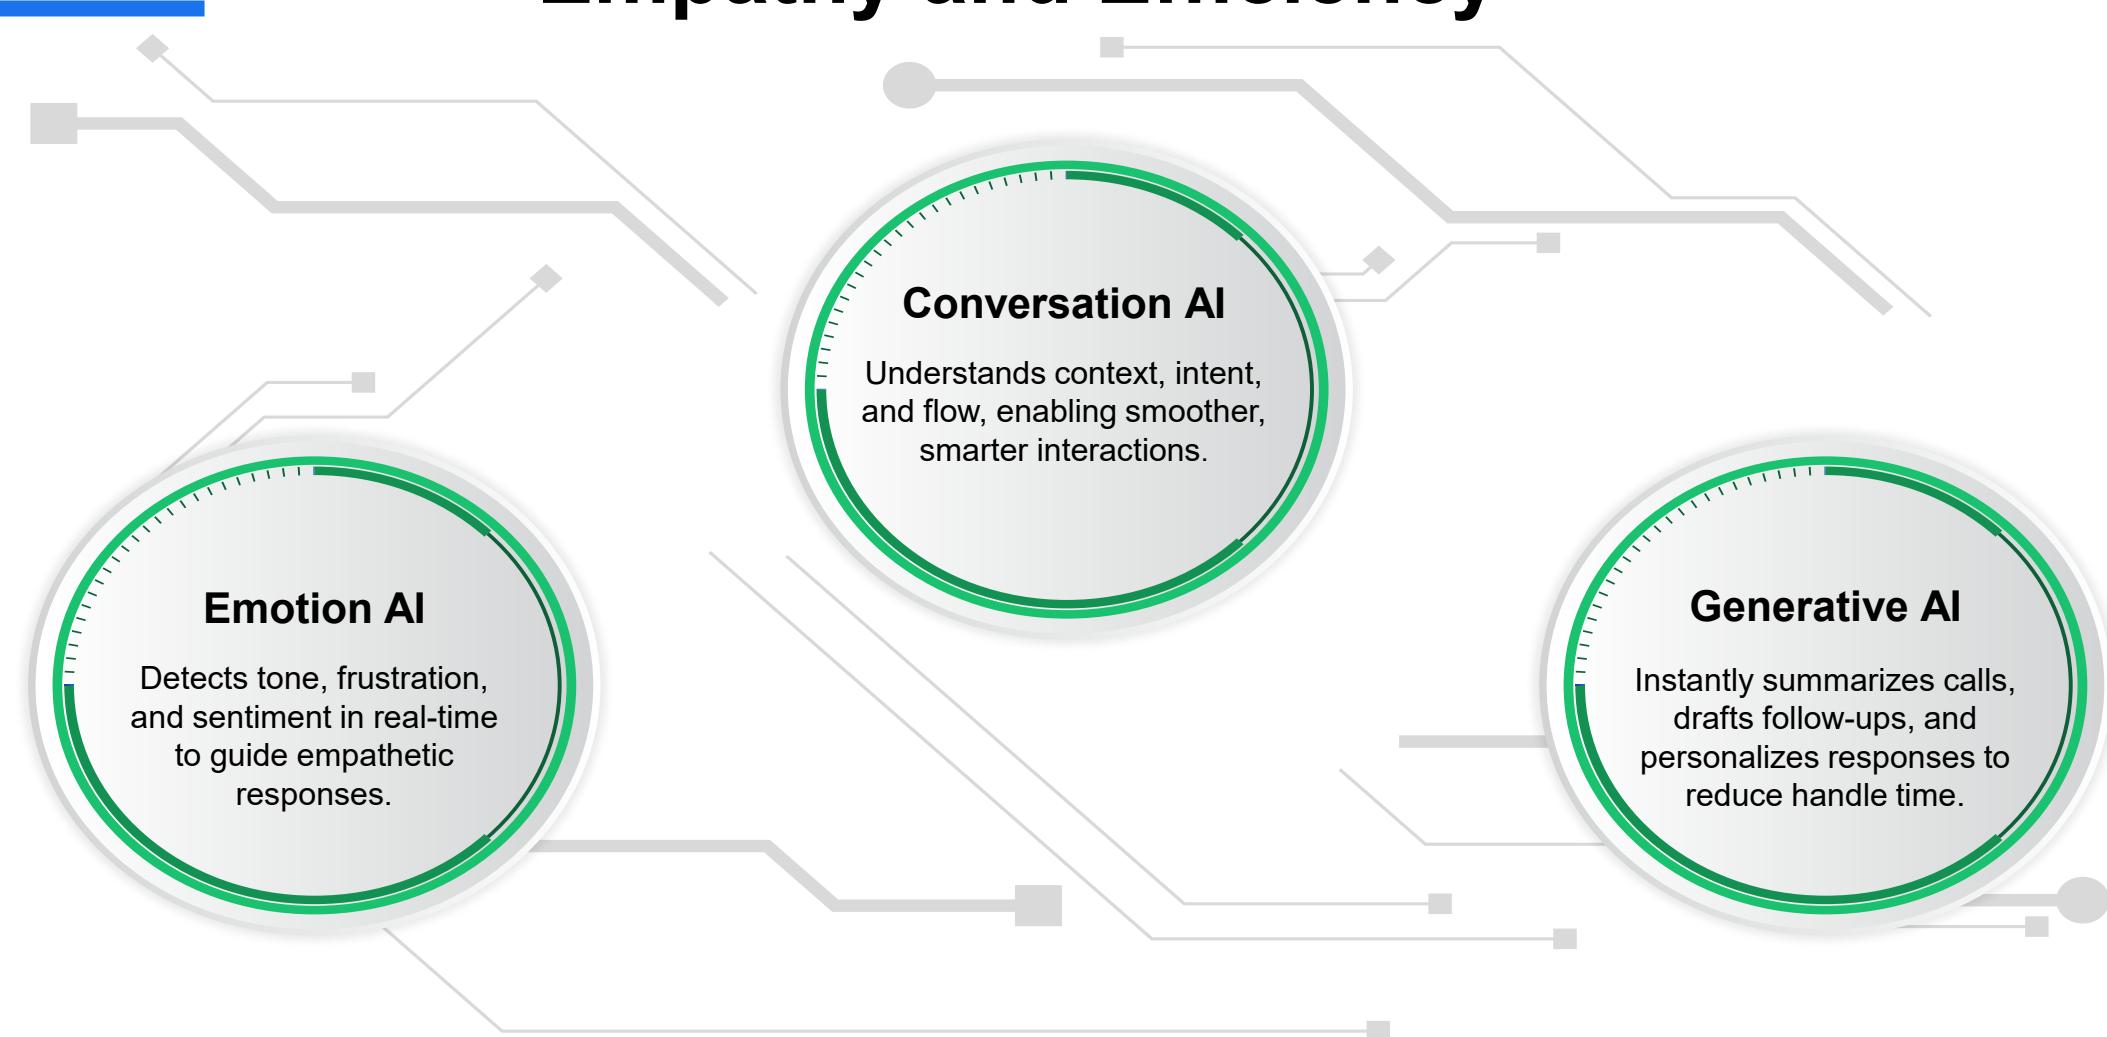
Product Vision



zoom

*To empower businesses
with AI to give people
immediate and nuanced
assistance so no one ever
has to be put on hold.*

Empowering Every Interaction With Intelligence, Empathy and Efficiency



By focusing on these three pillars, we aim to make communication frictionless and more human through Agentic AI

zoom

Research Methods

Primary Research Interviews

- 4 enterprise
- 4 end-user

Survey with end-users

N = 26

Secondary Research

Market Overview, Competitive analysis, pricing

Zach's Guidance

Focus and direct efforts



MARKET AND CUSTOMER UNDERSERVED NEEDS

Market Overview

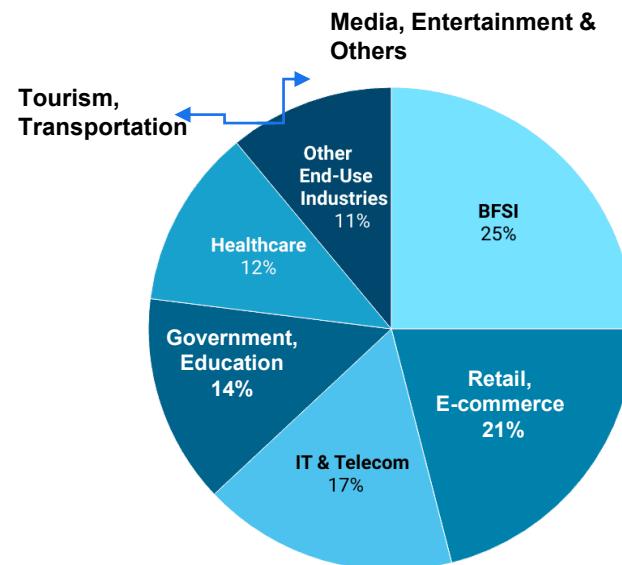


Source: Zoom CX & Virtual Assistant's Market Information



Source: [Precedence Research-CC](#),
[Precedence Research-VA](#)

Global Contact Center as a Service Market Share – By End-use Industry



Source: <https://scoop.market.us/contact-center-as-a-service-statistics/>



**Large enterprises accounted for 60%,
Size: 1,000+ tickets/day. (>30 agents)**
=> Market value: \$5.8B



**SME accounted for 40%
Size: 50 to 1,000 tickets per day (<30 agents)**
=> Market value: \$3.8B

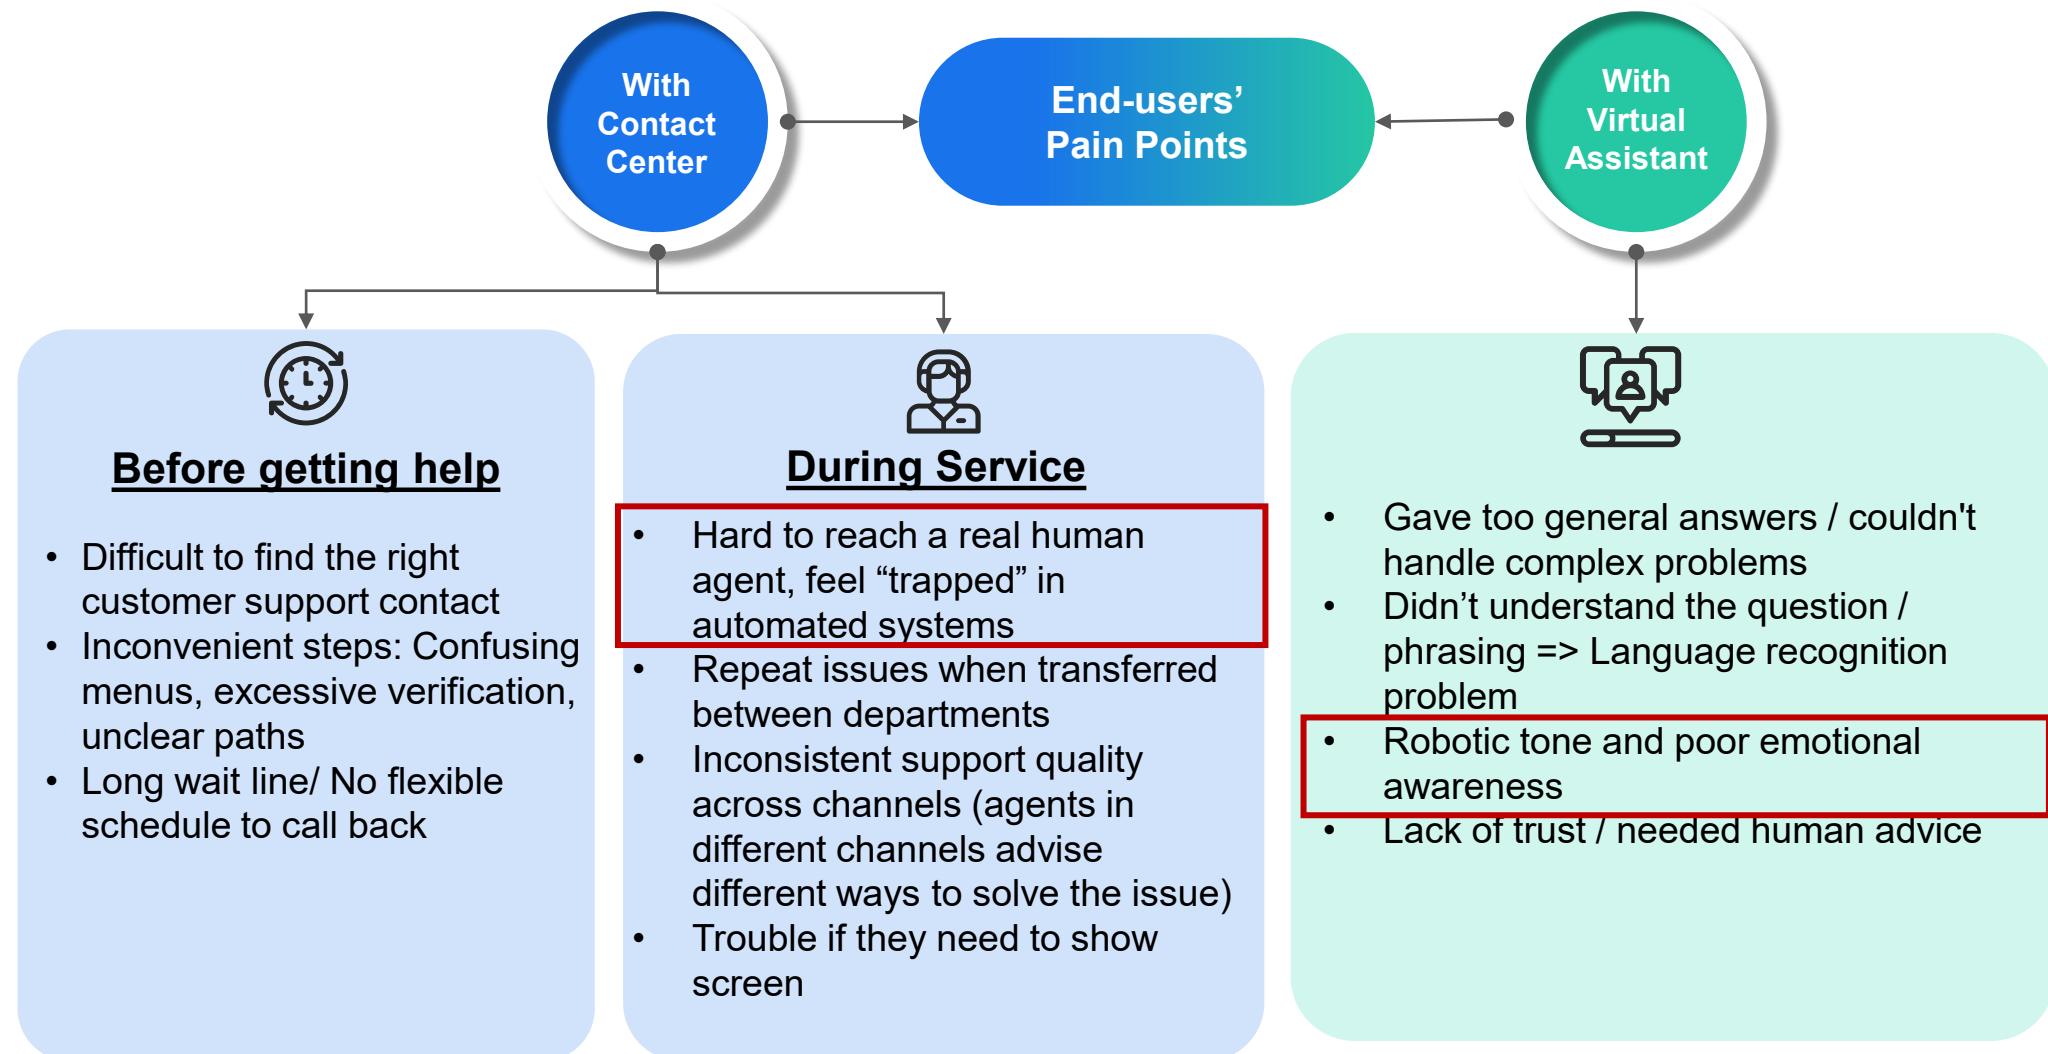
Source: [snsinsider.com](#) & Interview with Businesses Mar'25 by Team 2

Businesses' Underserved Needs

Source: Interview with Businesses in Mar'25 – Team 2

	Retail Scale: Vietnam SME: 300 tickets/ day Inhouse CC	Education Scale: US Large: 1000+ tickets/ day Inhouse CC	Healthcare Scale: Global SME & Large: 100 – 1000+ tickets/ day Outsourced CC	Telecom Scale: US Large: 10,000+ tickets/ day Outsourced CC
1. Feedback Sentiment & Insights Analysis Difficult to integrate customers' feedbacks from sources (text, call, email, chat, social media...), lack the ability to categorize issues & extract sentiment & insights to guide decision-making.	✓		✓	
2. Language and Communication Barriers Virtual agents struggle to understand and respond in human-like language, esp. when customers use informal expressions or need empathy.	✓	✓	✓	✓
3. System, Integration, and Tech Limitations Painful and time-consuming set-up process, inflexible integration with CRM platforms and backend systems, esp. when system delays or bugs interrupt live support and reduce first-contact resolution.	✓	✓	✓	
4. Staff, Scaling & Real-time support Limited staffing, lack of 24/7 service, & high agent turnover causes delayed responses, low resolution time and reputational risks, making it hard to scale up when the company meet high ticket volumes		✓	✓	✓
5. No clear SOP about products & support policies Lack of clear SOP about the products & customer support policy, leading to customers' confusion, weak cross-functional collaboration, repeated inquiries, & longer resolution times.	✓	✓	✓	✓

End-Users' Underserved Needs



Target Customers

Mass-Consumers

(High-volume, general public-facing)

- Language: Informal, conversational, incl. slang and abbreviations
- Emotion: Friendly, Quick, Polite

Specific-Consumers

(Sensitive, regulated, high-stakes support)

- Language: Formal, technical, precise vocab
- Emotion: Calm, patient, empathetic, expertise

Start with **US companies** -> expand globally

Start with **medium companies** with MVP, expand to large enterprises with complete products

Businesses

• By Industry:

- Banking, Financial Services & Insurance
- Retail & E-Commerce;
- IT & Telecom

Healthcare

Education & Government

• By Market:

US & Global

• By Size:

- Medium companies: 100 - 1,000 tickets per day (<30 agents)
- Large Enterprises: 1,000+ tickets/day. (>30 agents)

Prioritized Needs

- Language & Communication Barriers (language and emotion understanding)
- No clear SOP about products & support policies (how the end-user should be categorized to meet real agent)



COMPETITIVE LANDSCAPE VS ZVA



Zoom Virtual Agent: SWOT Snapshot in the CX AI Market

The contact center AI market is growing fast, with players like Genesys and Salesforce offering robust but complex, while startups like Ada and Sierra focus on automation. Zoom Virtual Agent stands out by combining emotional intelligence, fast deployment, and seamless Zoom integration, offering a more human and agile customer experience.

Your weaknesses	Limited native CRM integration and relatively low brand recognition in the CX space compared to Salesforce or Genesys.
Your strengths	Seamless integration with Zoom ecosystem, fast low-code deployment, RAG-powered generative engine, strong multichannel capabilities.
Competitor weaknesses	Sierra and Ada lack emotional AI depth; Genesys has a steep learning curve; Salesforce is expensive and complex to customize.
Competitor strengths	Salesforce excels in CRM-AI synergy; Genesys offers end-to-end CX automation; Ada is scalable; Sierra provides empathetic voice AI.
Your opportunities for growth/competitive advantage	Expansion into SMBs and Mid-Market, Autonomous smart control and NLP Advancements for human-like, emotion-aware responses.

Value Proposition

“Zoom Virtual Agent delivers fast, emotionally intelligent customer support through adaptable, human-like agent personas, low-code deployment and seamless integration.”

FEATURES AND ROADMAP



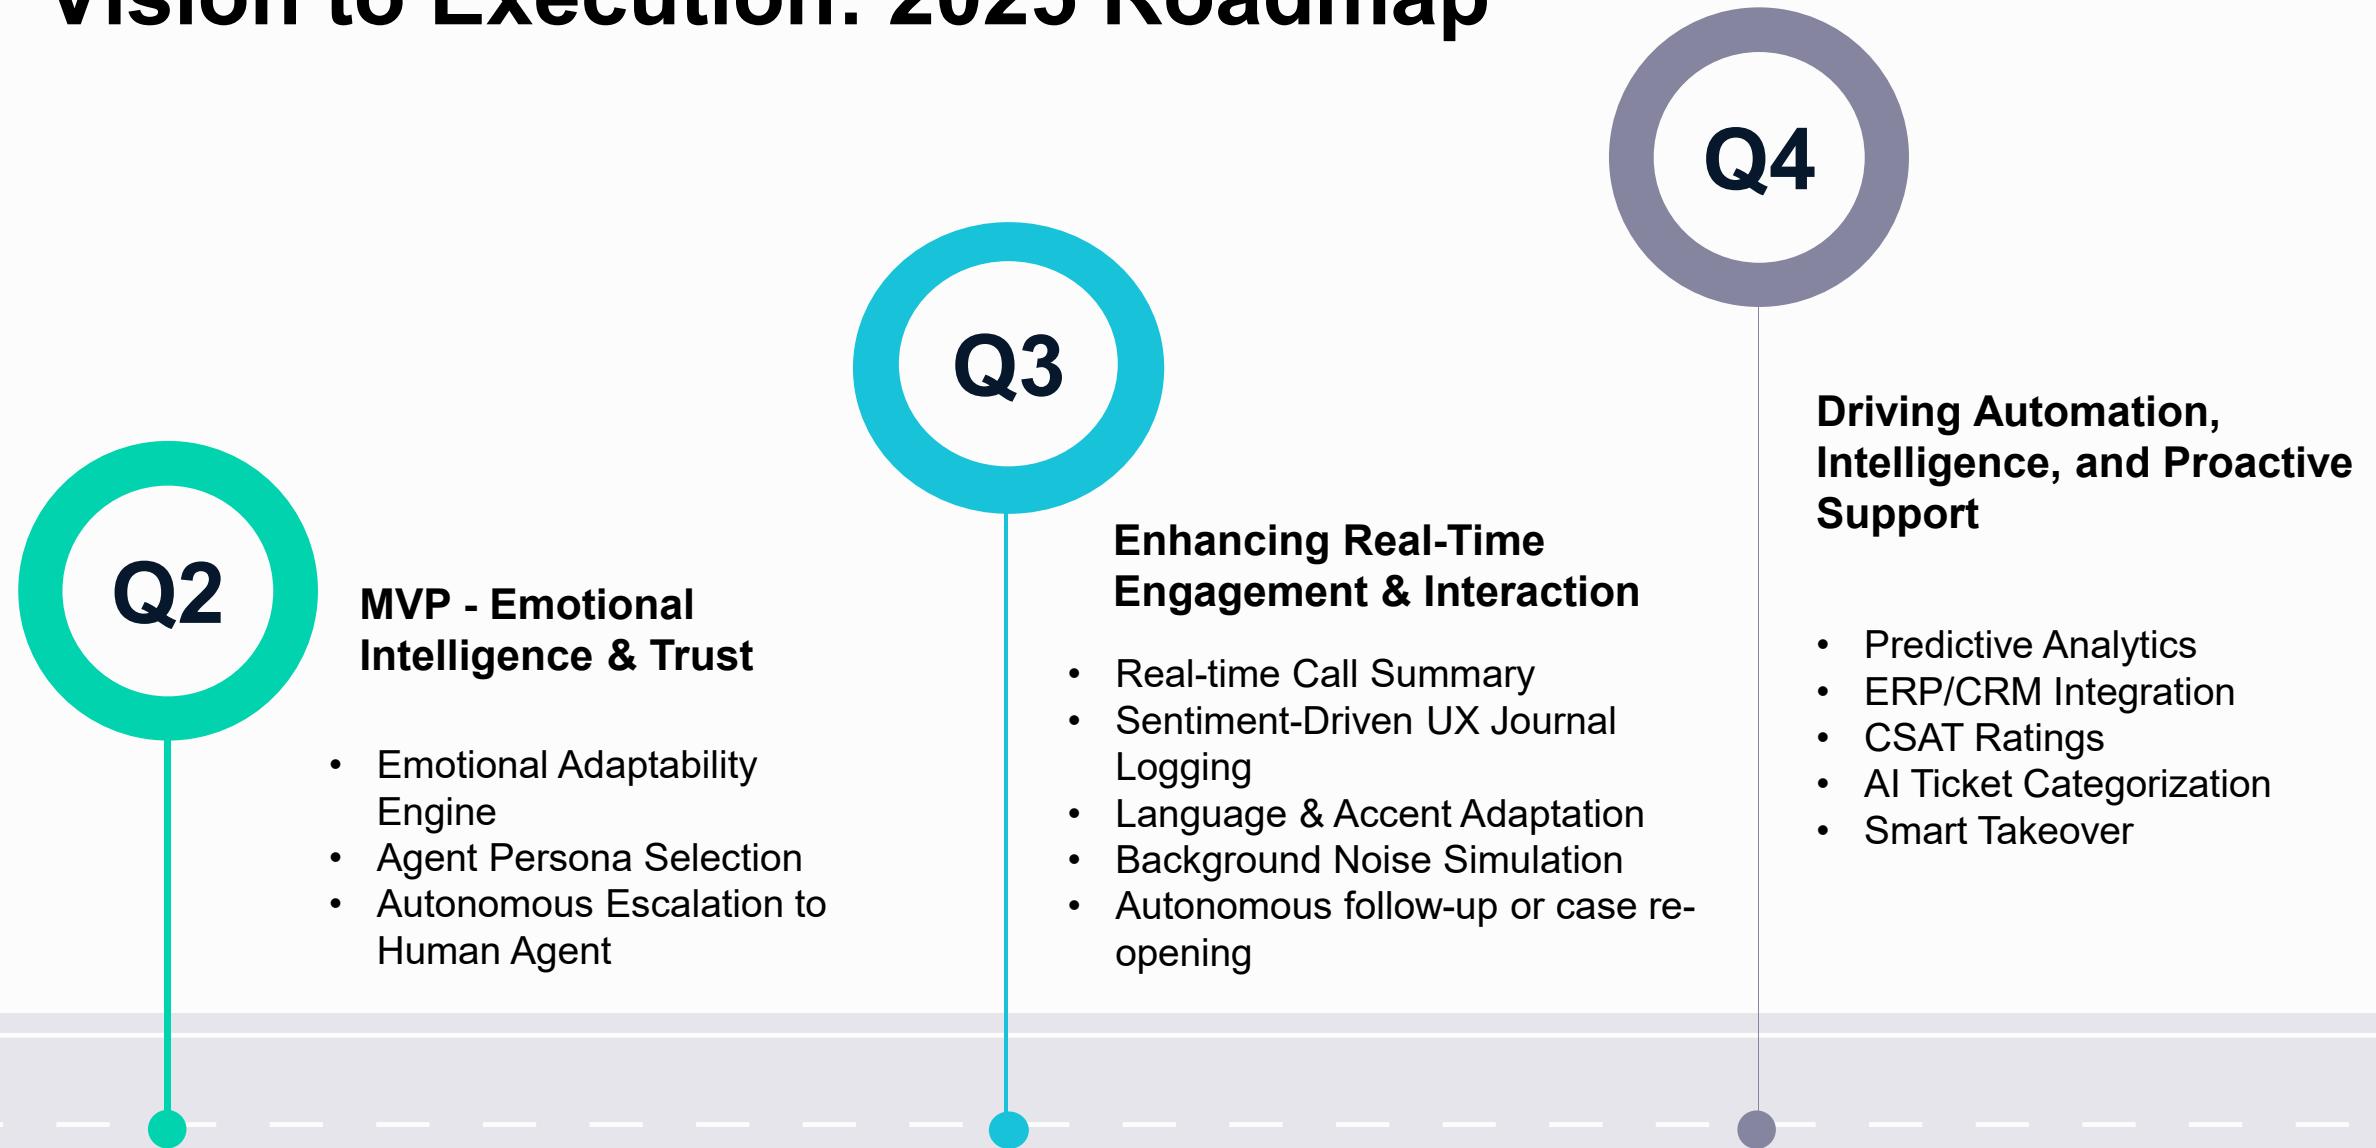
zoom

Feature Prioritization

Delighters	Performance Features	Basic Features	Indifferent
Autonomous follow-up or case re-opening	Emotional/frustrated keywords/tones detection	Ease of implementation & ERP/CRM integration	
Smart Takeover	Predictive analytics	Call and Text options	
Subtle background noise (to mimic call centers)	Less than 1-minute max wait time (fast GenAI processing) – Hold screen	Compliant categorization for ticketing	
User Experience Journal	Real-time Call Summary	CSAT Ratings	
Personality selections (mimic human agent)		Ticketing History	
Language and English accent selection		Autonomous Escalation to Human Agent	

Mapping features by what users need vs. what makes them stay, using the Kano Model.

Vision to Execution: 2025 Roadmap

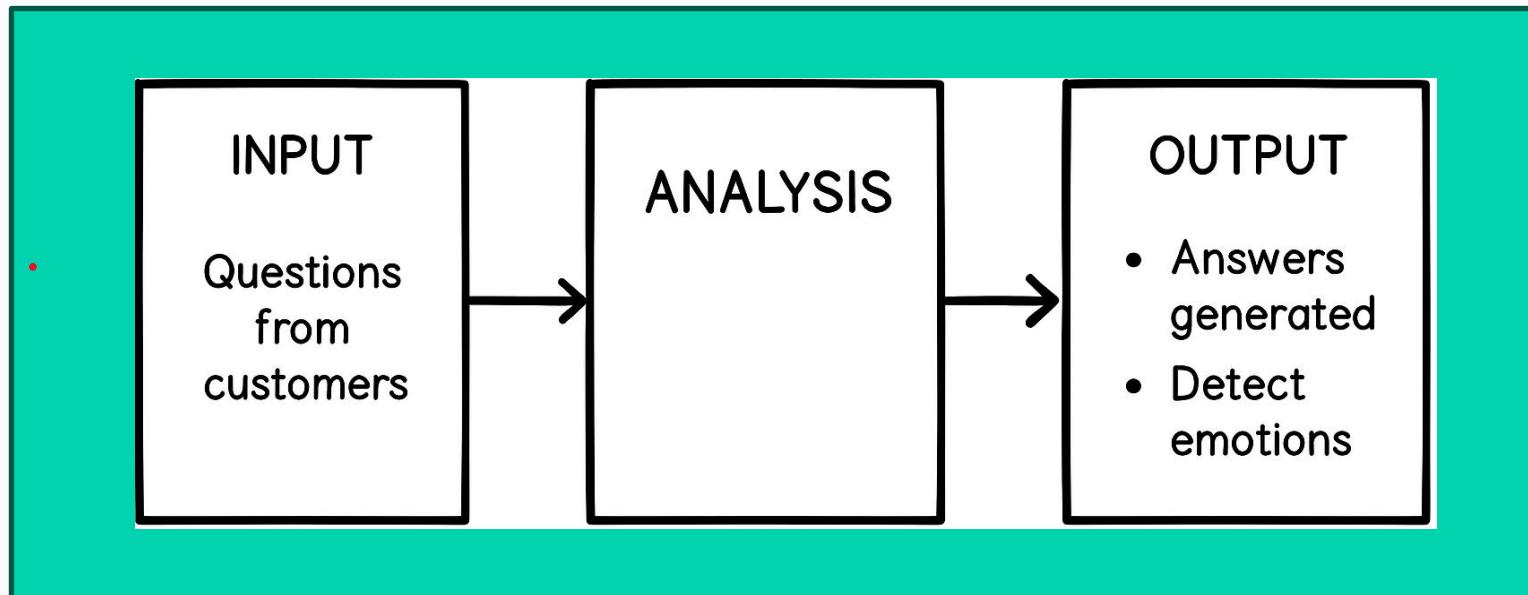


PROTOTYPES AND DEMOS



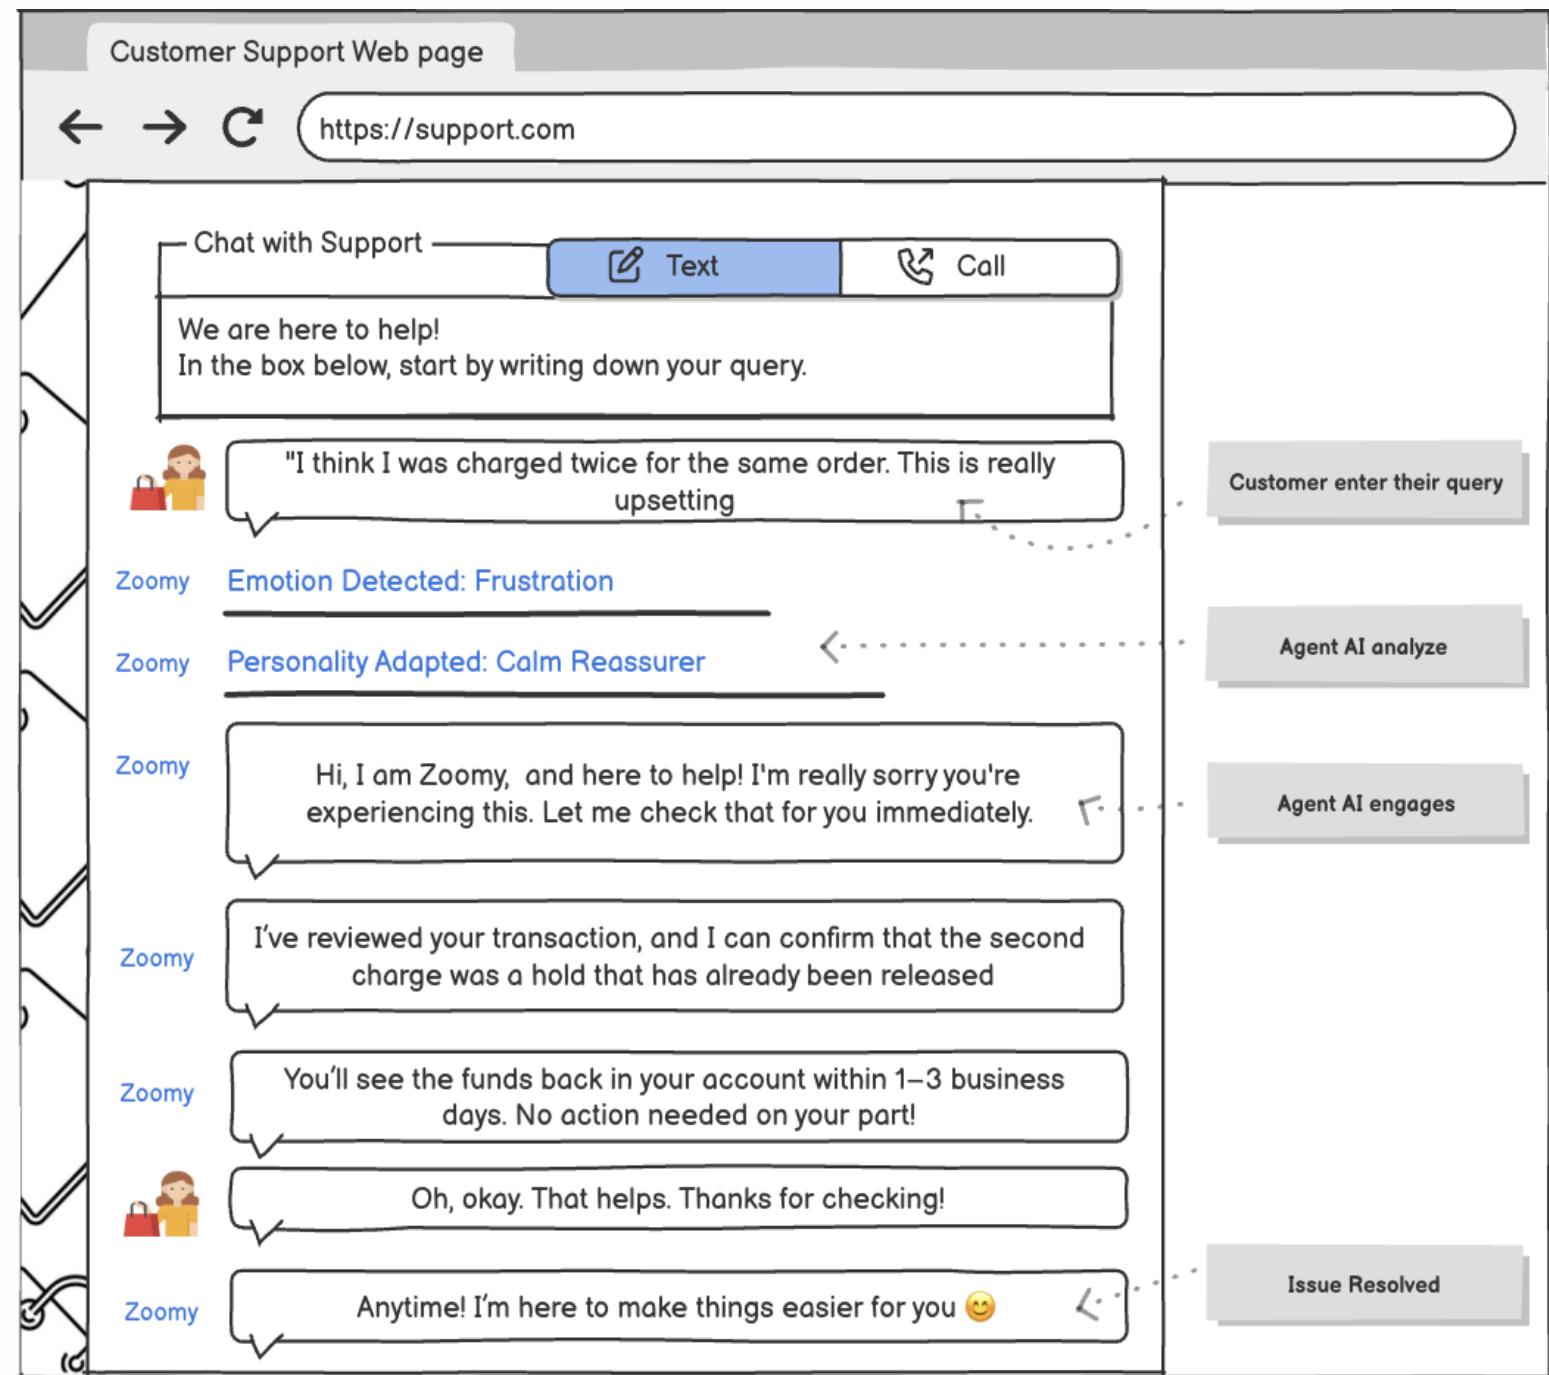
zoom

Solution Map



Detected Keywords	Interpreted Emotion	Personality Triggered
frustrated, upset, angry, annoyed	Frustration	Calm Reassurer
confused, overwhelmed, don't get it	Anxiety / Confusion	
worried, concerned, stressed	Concern / Stress	

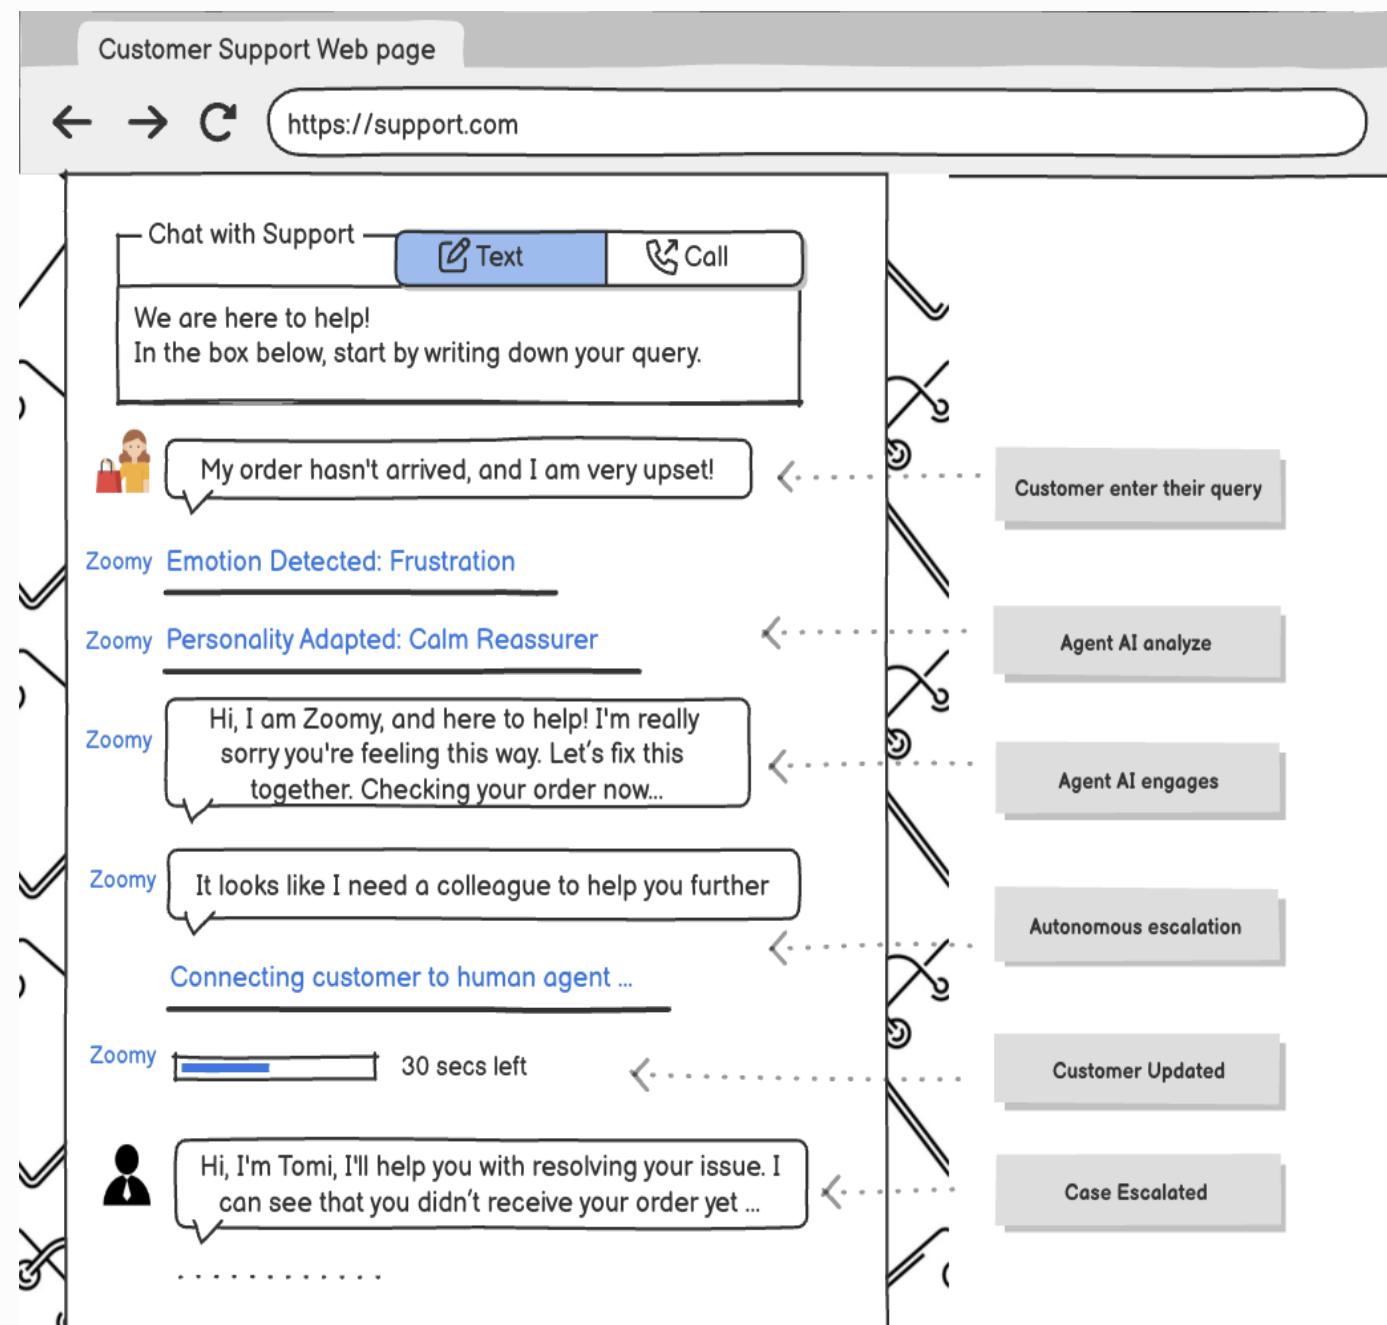
Emotionally Adaptive Resolution



Recognizing Limits and Triggering Escalation

Escalation Rules (examples):

- Agent can't find the customer data.
- Timed interaction.
- Repetition of "I need to talk to a customer representative similar phrases.



MVP Testing

Pass

Pivot

Persevere

As testing is still ahead, our decision in the meantime is based on strong user interest, pain points, and aligning with real-world customer service needs.

PRICING AND GTM STRATEGY



zoom

Pricing



Q2

- Industry: Healthcare, Education and Government
- Features: Emotional Adaptability, Autonomous Escalation
- Modality: Just text
- **Pricing: Per month: \$1100 ; \$.018 per interaction**



Q3

- Industry: Added Retail - Ecommerce, Telecom & IT, Financial Services
- Features: Emotional Adaptability and Autonomous Escalation for voice
- Pricing:
 - + **Education and Government - \$1100/month, \$.2/min, \$.018/interaction**
 - + **Retail - Ecommerce, Telecom & IT, Financial Services:**
 - > **Medium(up to 1000 tickets/day) - \$1500/month, \$.26/min**
 - > **Large(more than 1000 tickets/day) - \$2500/month (more personas selection for agents), \$.24/min**
- Feature Pricing: Real Time Call Summary : **\$890/month**

GTM Strategy

Target Segment	Small & Medium Enterprises in the US, in Healthcare, Education & Government industries	Medium & Large Enterprises, Global, Adding Retail, Telecom, BFSI industries		
MKT Strategy	Q2: Emotional Intelligence + Escalation	Q3: Real-Time Engagement & Interaction	Q4: Automation, Intelligence, Proactive Support	
Distribution Strategy	Channels	Website & Blogs, LinkedIn	Q2 + PR, Thought Leader Endorsements, Product Launch Events/Tutorial Webinars , Social Media (FB, YT, IG, TikTok), SEO/SEM, In-app banners + notifications, EDM	Q3 + Advanced social ads, Account-based email marketing, Industry CX events
KPIs	Activities & Content	Demo + Tutorial Video/Article/Post	Q2 + Product Showcase/Customer testimonials/ Case studies featuring content White paper on Agentic AI in CX; Launching events/ Panel at CX event (e.g., CCW, CXNext) ROI calculator tool (customers can estimate impacts when applying our system)	
Direct Sales (Account-Based Distribution)	Pilot with 3 key accounts, Outreach to Zoom small-business customers for free trials	Add to Zoom & CRM marketplaces, Expand to medium accounts	Full pipeline push targeting large enterprises, Scale sales team	
Strategic Partnerships		Co-market with voice AI (e.g., BehavioralSignals)	Co-market with CRM vendors (Salesforce, Zendesk) with cross-sell campaigns	
	3 pilots running demo to 50 customers, 10 trial customers 30% increase CSAT on average per company	100 new medium clients 60% increase CSAT on average per company	50+ new large enterprise clients 100% increase CSAT on average per company	

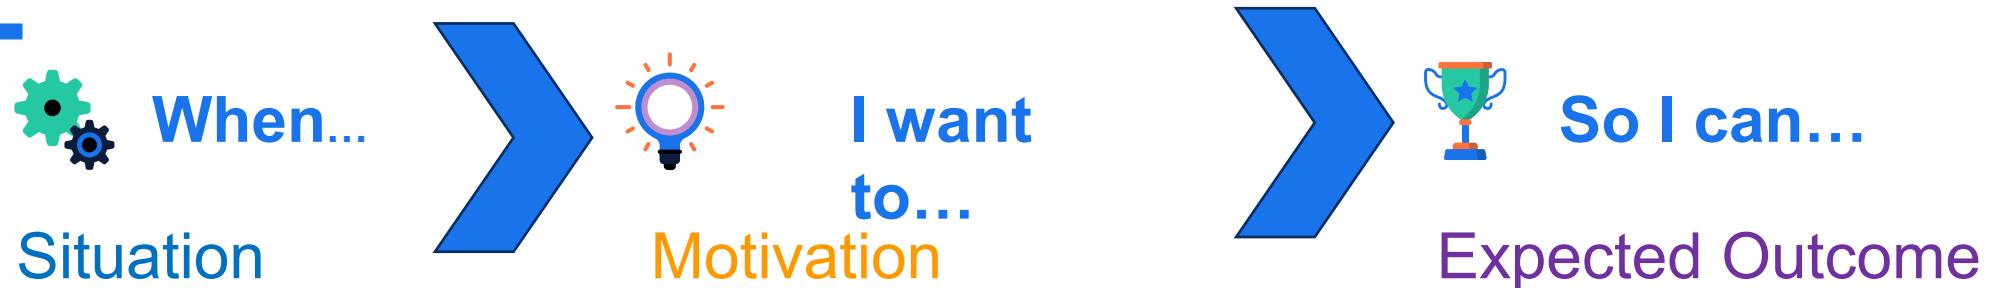
Thank You!



zoom

Appendix

JTBD



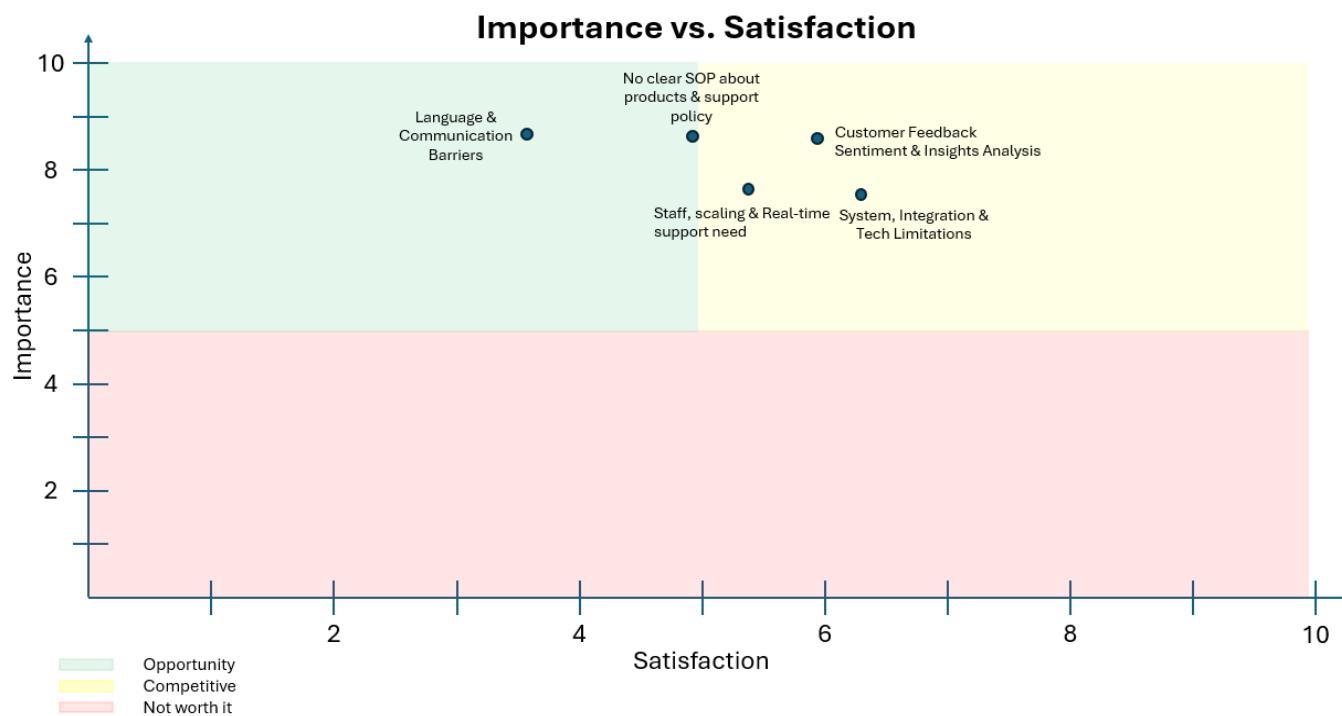
- When customer support teams are overwhelmed, I want to automate routine inquiries so they can focus on complex, high-value cases.
- When a customer reaches out, I want the system to understand emotional tone and urgency so the response can be personalized and escalated appropriately
- When launching a virtual agent, IT teams want to integrate with existing CRMs and ERPs quickly so they can go live without major development delays.
- When customers call from diverse regions, businesses want to offer preferred languages and accents to make interactions feel natural and localized.
- When a call ends, managers want to access real-time summaries so they can track performance and follow up efficiently.

User Stories

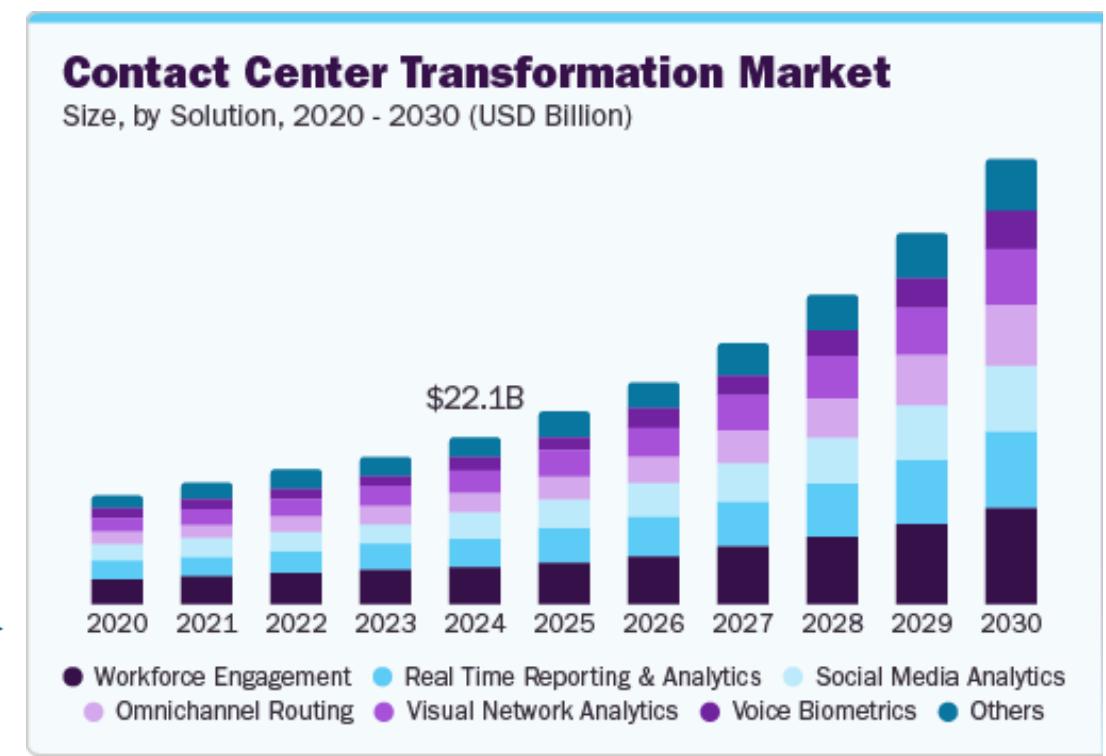
Feature	User Story	Acceptance Criteria
Language and English Accent Selection	As an admin, I want to configure available languages and regional English accents for the AI assistant, So that our customers can interact in a way that feels natural and regionally appropriate.(e.g., US, UK, Indian, Australian)	<ul style="list-style-type: none">Admin can enable/disable specific languages and accents via a configuration panel.System supports regional variants (e.g., US, UK, Australian, Indian English).Selected accents persist across sessions unless changed.Voice output is regionally accurate and tested for clarity.Optionally logs user language/accent preferences for future CX personalization.
Personality Selections	As an admin, I want to set the AI assistant's personality style or allow users to select it, so that the interaction tone aligns with our brand and customer expectations.(e.g., friendly, formal, humorous)	<ul style="list-style-type: none">Admin can set a default AI personality (e.g., formal, casual, enthusiastic).Admin can allow or restrict personality selection by users.Personality style affects AI's tone, phrasing, and pace.Personality choices are logged for analytics or A/B testing.Admin can preview and test personality tones during setup.
Real-time Call Summary	As an admin, I want the AI to generate real-time call summaries and attach them to tickets, so that agents and supervisors can quickly review key takeaways without manual notes.	<ul style="list-style-type: none">AI generates structured summaries immediately after each interaction.Summary includes issue type, customer intent, resolution status, and next steps.Admins can configure which summary fields are mandatory.Summaries are attached to CRM/ticketing system automatically.Agents can edit and approve summaries before final submission.
Emotional/Frustrated Keywords/Tones Detection	As an admin, I want the AI to detect emotional or frustrated keywords and tones during conversations, so that agents and supervisors can intervene or adjust strategy to improve customer experience.	<ul style="list-style-type: none">AI detects vocal tone and keyword patterns indicating frustration or emotional distressAdmin can set sensitivity thresholds and customize keyword listsReal-time alerts sent to supervisors or flagged to agents mid-callEmotional flags logged and linked to call records for reviewDashboard includes sentiment trend analytics and agent-level insightsOptional escalation rules triggered by repeated emotional incidents

Businesses' Underserved Needs vs Available Solutions in Market

Needs	Available Solution in Market
Language & Communication Barrier	Voice Biometrics
No clear SOP + Staff Shortage	Workforce Engagement/Omnichannel Routing
Customers' Feedback Sentiment & Insight Analysis	Realtime Reporting/Social Media/Visual Network Analytics



Source: Interview with Businesses in Mar'25 – Team 2



Current Landscape: How Zoom Virtual Agent Compares

This matrix outlines how Zoom Virtual Agent's current features stack up against leading competitors across five core functional benefits, highlighting strengths and potential areas for growth.

	Emotional Intelligence	Conversation Quality	Generative Capabilities	Omnichannel Coverage	Customization & Persona Control
Zoom Virtual Agent	Basic sentiment detection via NLP	Fast, accurate resolution	RAG-based response engine	Full CRM omnichannel (email, chat, social)	Basic admin tools, limited persona UI
Salesforce Agentforce	Deep LLM-based sentiment analysis	Autonomous agents with CRM memory	Task execution with GenAI	Full CRM omnichannel (email, chat, social)	Deep customization via Salesforce Platform
Genesys Virtual Agent	Stress & sentiment analysis	Agent assist + contextual recall	Auto-generated intents & summaries	Voice, chat, mobile	Custom flows & conversation maps
Ada CX	Limited emotional awareness	Context-aware, multi-turn responses	GenAI-powered answers + automation	WhatsApp, chat, email, SMS	Rule-based workflows
Sierra AI	Empathy-tuned voice AI	Human-like voice responses	Voice-personalized GenAI agents	Voice-first + multi-channel	Voice tone & persona builder

MVP Launch: Emotionally Adaptive AI

Theme: Build empathy and trust in AI support

MVP Objective: Deliver the minimum lovable product with emotional intelligence, personality, and human fallback.

Feature	Description
Emotional Adaptability	Detects emotional cues like frustration and adjusts tone of response.
Personality Selection (1 Option)	Allows users to choose a communication style to personalize interactions.
Autonomous Escalation to Human	Triggers human takeover when the AI reaches capability or emotional limits.

Pricing Info(Appendix)

Q2

- Education and Government
- Costs:IBM Watson Tone analyzer - .003NLU for first 1-250000 NLU items, .001 up to 5,000,000, .0002 above
- NLU = natural language understanding 1 conversation avg 6NLU =.018/interaction
- Humantic AI – personality match \$1000/month
- Benefit/Value - Reduction in labor, increase in retention rate in clv
- 5% increase in retention for every 1% increase in satisfaction = 1% increase in revenue
- Help 24/7 - Reduction in calls – fewer staff \$18/hour

Q3

- Costs – behavioral signals .012/min,
- Speech to text .024/min Google, .22 eleven labs
- \$1100/month, .05/1000 addl char over 22M eleven labs
- Reduction in calls to humans
- More substantial increase in satisfaction
- Per month: \$1100, $(.22+.012) * \min(\text{tier for industry}) + 2\% \text{ avg revenue for industry} + \text{size(tier)} + \text{size} * 50\% * 18\$$
- Cost 600\$/month - dragoncopilot

Cost Elements Info(Appendix)

Tool 1: IBM Watson Tone Analyzer => Emotion Detection For Text

Example:

- Each chat has **12 messages** (6 from customer, 6 from agent/AI)
- You want to analyze only **customer messages** (6 per conversation)

Then:

• **1 conversation ≈ 6 NLU items**

If you process **10,000 conversations/month**, you'll use:

• **$10,000 \times 6 = 60,000 \text{ NLU items/month}$**

With IBM NLU Tier 1 pricing:

• **$60,000 \times \$0.003 = \$180/\text{month}$**

https://cloud.ibm.com/docs/natural-language-understanding?topic=natural-language-understanding-pricing&utm_source=chatgpt.com

Lite

The Lite plan is free for users who are looking to try out Natural Language Understanding or build a proof-of-concept.

Pricing

- Free 30,000 NLU items/month

Custom Model

- 1 free custom model

Standard

A "pay-as-you-go" plan that is recommended once you are ready to move your application from proof-of-concept to production.

Pricing (billed monthly)

- Tier 1: \$0.003/NLU item for the first 1-250,000 NLU items
- Tier 2: \$0.001/NLU item for 250,001 to 5,000,000 NLU items
- Tier 3: \$0.0002/NLU item for additional NLU items past 5,000,000

Custom Entities and Relations Models (\$/custom model/month)

- \$800 for all tiers

Custom Classification Model (\$/custom classification model/month)

- \$25 for all tiers

Tool 2: Behavioral Signals => Emotion Detection For Voice

Standard Plan

- **Batch Processing:** \$0.012 per minute.
- **Streaming Processing:** \$0.015 per minute.
- **Features:**

- No long-term contracts.
- Cloud deployment.
- Online support.
- Data retention for 5 years.

Bulk Discounts

For larger volumes, discounts are available:

- **\$500–\$999:** 5% bonus credits.
- **\$1,000 and above:** 10% bonus credits.

Custom Plan

Tailored solutions for specific requirements:

- **Features:**
- Volume-based discounts.
- Custom builds.
- Tailored models.
- On-premises deployment.
- Dedicated success manager.
- Premium support.

**OLD PRICING STRUCTURE => NOW
THEY JUST DO CUSTOMIZED PRICING**

<https://behavioralsignals.com/home>

Tool 3: Tool to generate the answers and adapt emotions for the agents

•OpenAI – Natural Language Processing (NLP)

Service	Purpose	Model / Tool	Estimated Price	Section on Pricing Page													
Chat Completions	Generate replies, rewrite tone based on emotion	gpt-4 (uses GPT-4 Turbo)	\$0.01 / 1K input tokens \$0.03 / 1K output tokens	Our APIs	<ul style="list-style-type: none">◆ What is a “token”? <p>A token is a chunk of text — usually a word or part of a word.</p> <p>Here's how it works in practice:</p> <table><thead><tr><th>Text</th><th>Token Count</th></tr></thead><tbody><tr><td>Hello</td><td>1 token</td></tr><tr><td>ChatGPT is awesome!</td><td>5 tokens</td></tr><tr><td>I'd like to speak to a human.</td><td>8 tokens</td></tr><tr><td>antidisestablishmentarianism</td><td>1 token</td></tr><tr><td>Zoom AI:</td><td>2 tokens</td></tr></tbody></table>	Text	Token Count	Hello	1 token	ChatGPT is awesome!	5 tokens	I'd like to speak to a human.	8 tokens	antidisestablishmentarianism	1 token	Zoom AI:	2 tokens
Text	Token Count																
Hello	1 token																
ChatGPT is awesome!	5 tokens																
I'd like to speak to a human.	8 tokens																
antidisestablishmentarianism	1 token																
Zoom AI:	2 tokens																
Embeddings	Match past chats, FAQ search, personalization	text-embedding-3-small	\$0.00002 / 1K tokens	Our APIs													
Text-to-Speech (TTS)	Convert AI replies to voice (if needed for voice support bots)	tts-1, tts-1-hd	\$0.015–\$0.03 / 1K characters	Text-to-speech (TTS)	<ul style="list-style-type: none">◆ Rule of Thumb:<ul style="list-style-type: none">• 1,000 tokens ≈ 750 words• Average English sentence = ~15–25 tokens												
Moderation	Ensure AI responses stay safe and appropriate	Moderation endpoint	Free	Moderation	<p>So when OpenAI says:</p> <p>\$0.01 per 1,000 input tokens, \$0.03 per 1,000 output tokens (for gpt-4),</p>												
Assistant API (Optional)	If you want persistent memory, tools, or function calling	gpt-4 or gpt-3.5	Same as chat pricing (GPT-4 Turbo)	Assistants API (optional)													

Source: <https://openai.com/api/pricing/>

Tool 4: Voice AI Agent - Google ASR: Speech to Text

Cloud Speech-to-Text pricing

Cloud Speech-to-Text On-Prem is priced based on the amount of audio successfully processed by the service **each month**, measured in increments rounded up to 15 seconds.

You can view your current billing status, including usage and your current bill, in the [Cloud console](#). For more details about managing your account, see the [Cloud billing documentation](#) or [billing and payments support](#).

Pricing Table

Feature	Standard models	
	0-60 Minutes	Over 60 Mins up to 1 Million Mins
Speech Recognition	Free	\$0.006 / 15 seconds ** 0.024\$/min

*** Each request is rounded up to the nearest increment of 15 seconds.*

Pricing calculations

Each request is rounded up to the nearest increment of 15 seconds. For example, if you make three separate requests, each containing 7 seconds of audio, you are billed \$0.018 USD for 45 seconds (3×15 seconds) of audio. Fractions of seconds are included when rounding up to the nearest increment of 15 seconds. That is, 15.14 seconds are rounded up and billed as 30 seconds.

Source: <https://cloud.google.com/speech-to-text/priv/pricing>

Tool 5: Voice AI Agent – ElevenLabs

- Text to Speech
- Pricing: Business Plan: \$1,100/month for 22 mil characters

Source: <https://elevenlabs.io/pricing>

Models	Model <small>i</small>	Display In						Billing
	MULTILINGUAL	FLASH	CHARACTERS	MINUTES	MONTHLY	ANNUAL		
 Text to Speech								
 Speech to Text		Free	Starter	Creator	Pro	Scale	Business	
 Voice Changer		\$0/mo	\$4.17/mo \$5 <small>i</small>	\$18.33/mo \$22 <small>i</small>	\$82.5/mo \$99 <small>i</small>	\$275/mo \$330 <small>i</small>	\$1,100/mo \$1,320 <small>i</small>	
 Sound Effects		GET STARTED	GET STARTED	GET STARTED	GET STARTED	GET STARTED	GET STARTED	
 Voice Isolator		Characters included <small>i</small>	20,000	60,000	200,000	1,000,000	4,000,000	22,000,000
Products		Additional characters <small>i</small>	<small>x</small>	<small>x</small>	\$0.15/1000	\$0.12/1000	\$0.09/1000	\$0.05/1000
 Dubbing		Audio quality <small>i</small>	128 kbps, 44.1kHz	128 kbps, 44.1kHz	128 & 192 kbps (via API), 44.1kHz	128 & 192 kbps (via Studio & API), 44.1kHz	128 & 192 kbps (via Studio & API), 44.1kHz	
 Voices		<small>Free tier requires attribution and does not have commercial licensing</small>						
 Studio		SHOW API DETAILS <small>v</small>						

Tool 5: Voice AI Agent – ElevenLabs

- ElevenLabs – Speech to Text
- Pricing: Business Plan: \$1,100/month for 6000 hours

Models	Interface	Billing						
	Text to Speech	API	UI	MONTHLY	ANNUAL			
 Speech to Text		Free \$0/mo	Starter \$4.17/mo \$5 ⓘ	Creator \$18.33/mo \$22 ⓘ	Pro \$82.5/mo \$99 ⓘ	Scale \$275/mo \$330 ⓘ	Business \$1,100/mo \$1,320 ⓘ	
 Voice Changer		GET STARTED	GET STARTED	GET STARTED	GET STARTED	GET STARTED	GET STARTED	
 Sound Effects								
 Voice Isolator		Hours included ⓘ	2 hours 30 minutes	12 hours 30 minutes	62 hours 51 minutes	300 hours	1100 hours	6000 hours
Products								
 Conversational AI		Price per included hour	✖	\$0.4	\$0.35	\$0.33	\$0.3	\$0.22
 Dubbing		Price per additional hour	✖	✖	\$0.48	\$0.4	\$0.33	\$0.22
 Voices		Free tier requires attribution and does not have commercial licensing						
 Studio		0.003\$/min						